

Efficient Serial Programming with NumPy

Nick Featherstone
CU Research Computing

[Web Link to These Slides](#)

Performance & Python

Python is an **interpreted** language

- No binary executable is created
- Interpreter executes source code line-by-line
- Instructions are executed *naively*
 - exactly as you wrote them
 - In the order you wrote them
 - no inherent vectorization

Compiled languages are different

- A binary executable is created
- Instructions are executed *holistically*
 - Same outcome as naive approach
 - Compiler-assisted optimization & vectorization
 - Implementation differs between compilers

Python with Numpy

NumPy provides some benefits of a compiled language within Python's interpreted framework

It offers

- Arrays
(efficient memory layout)
- Array methods
(vectorized loop operations)

$$A = [7, 2, 18, 3]$$

memory layout: lists



non-contiguous

memory layout: arrays



contiguous

The Big Picture

...if you remember nothing else...

Whenever Possible:

- Use NumPy arrays instead of lists
- Use in-place operations
- Use array syntax instead of explicit loops

Getting started with NumPy

- Open [initialization.py](#)
- Must import the NumPy module in order to use its features

```
import numpy
```

Common import patterns

OR

```
import numpy as np
```

We use this

[Open this file:](#)

Parallelization Workshop /
Day3-Parallel_Python /
session1_numpy /
examples /
[initialization.py](#)

NumPy Documentation:

<https://docs.scipy.org/doc/numpy/user/index.html>

NumPy ndarray Initialization Pattern

```
import numpy as np  
my_array = np.init_type (dims, dtype='data type' )
```

`init_type` : describes whether to initialize array to zero or not
 zeros -> initialize array with zero values
 empty -> do not initialize array values

`dims` : tuple with dimensions (nx, [ny, nz, ...]) of the array
 e.g., (10), (10,2), (2,8,10)

`dtype` : string variable describing the type of variable
 e.g., 'int16', 'int32', 'float16', 'float32', 'float64', 'complex64'

more on data types: <https://docs.scipy.org/doc/numpy/user/basics.types.html>

Initializing Arrays with Values

Initialize using values from list

```
list = [ 0, 2, 1, 3]  
my_array = np.array (list, dtype='data type' )
```

Initialize using values on [a,b) with integer spacing n

```
my_array = np.arange (a, b, n, dtype='data type' )
```

Initialize using n evenly space values on [a, b]

```
my_array = np.linspace (a, b, n, dtype='data type' )
```

Quick Exercises

- Create a 1-D NumPy array with 3 16-bit integer elements, initialized to 0.
- Create a 1-D NumPy array with 4 64-bit floating-point values initialized to [0, 0.1, 0.2, 0.3] using *linspace*
- Create a 1-D NumPy array with 4 64-bit floating-point values initialized to [1.0 , 0.1 , 9.5 , 11.0] using *array*

Use Arrays Instead of Lists

A calculation using NumPy arrays, in conjunction with array syntax, will usually be faster than one using lists.

[Open this file:](#)

Parallelization Workshop /
Day3-Parallel_Python /
session1_numpy /
examples /
[arrays_vs_lists.py](#)

Avoid Loops When Possible

```
a = np.linspace(...)
b = np.linspace(...)
c = np.zeros(...)
```

```
for i in range(n):
    c[i] = a[i]*b[i]
```

explicit loop
not vectorized



```
c = a*b
```

array syntax
vectorized!

[Open this file:](#)

Parallelization Workshop /
Day3-Parallel_Python /
session1_numpy /
examples /
[noloops.py](#)

Exercise1

Rewrite this program using

- NumPy arrays instead of lists
- array syntax instead of loops

[Open this file:](#)

Parallelization Workshop /
Day3-Parallel_Python /
session1_numpy /
exercises /
[exercise1.py](#)

In-Place Operations

When possible, use in-place operations to avoid unnecessary copies

- `a = a+2` -> `a += 2`
- `a = a-2` -> `a -= 2`
- `a = a*2` -> `a *= 2`
- `a = a/2` -> `a /= 2`

[Open this file:](#)

Parallelization Workshop /
Day3-Parallel_Python /
session1_numpy /
examples /
[inplace.py](#)


Array Ordering

- N-D arrays reside in 1-D Memory
- Two different ways of storing arrays

$$A = \begin{bmatrix} a_{00} & a_{01} \\ a_{10} & a_{11} \end{bmatrix}$$


Row-major: stripe row-by-row (C/C++; PYTHON DEFAULT)

Last index is “fastest”



Column-major: stripe column-by-column (Fortran)

First index is “fastest”



Array Ordering

- We can control the ordering if desired

[Open this file:](#)

Parallelization Workshop /
Day3-Parallel_Python / session1_numpy
/examples / [ordering.py](#)

Row-major: stripe row-by-row (C/C++; PYTHON DEFAULT)

Last index is “fastest”

a_{00}	a_{01}	a_{10}	a_{11}
----------	----------	----------	----------

Column-major: stripe column-by-column (Fortran)

First index is “fastest”

a_{00}	a_{10}	a_{01}	a_{11}
----------	----------	----------	----------

Array Ordering: Why Care?

- Sometimes, you REALLY do have to write a loop
- The innermost loop should correspond to the fastest array index

[Open this file:](#)

Parallelization Workshop /
Day3-Parallel_Python /
session1_numpy /
examples /
[access_patterns.py](#)

Row-Major

```
for i in range(m)  
    for j in range(n):  
        a+=b[i][j]
```

Column-Major

```
for j in range(n)  
    for i in range(m):  
        a+=b[i][j]
```

I/O with Numpy Arrays

- Writing/reading numpy arrays to/from a file is easy...

[Open this file:](#)

Parallelization Workshop /

Day3-Parallel_Python /

session1_numpy /

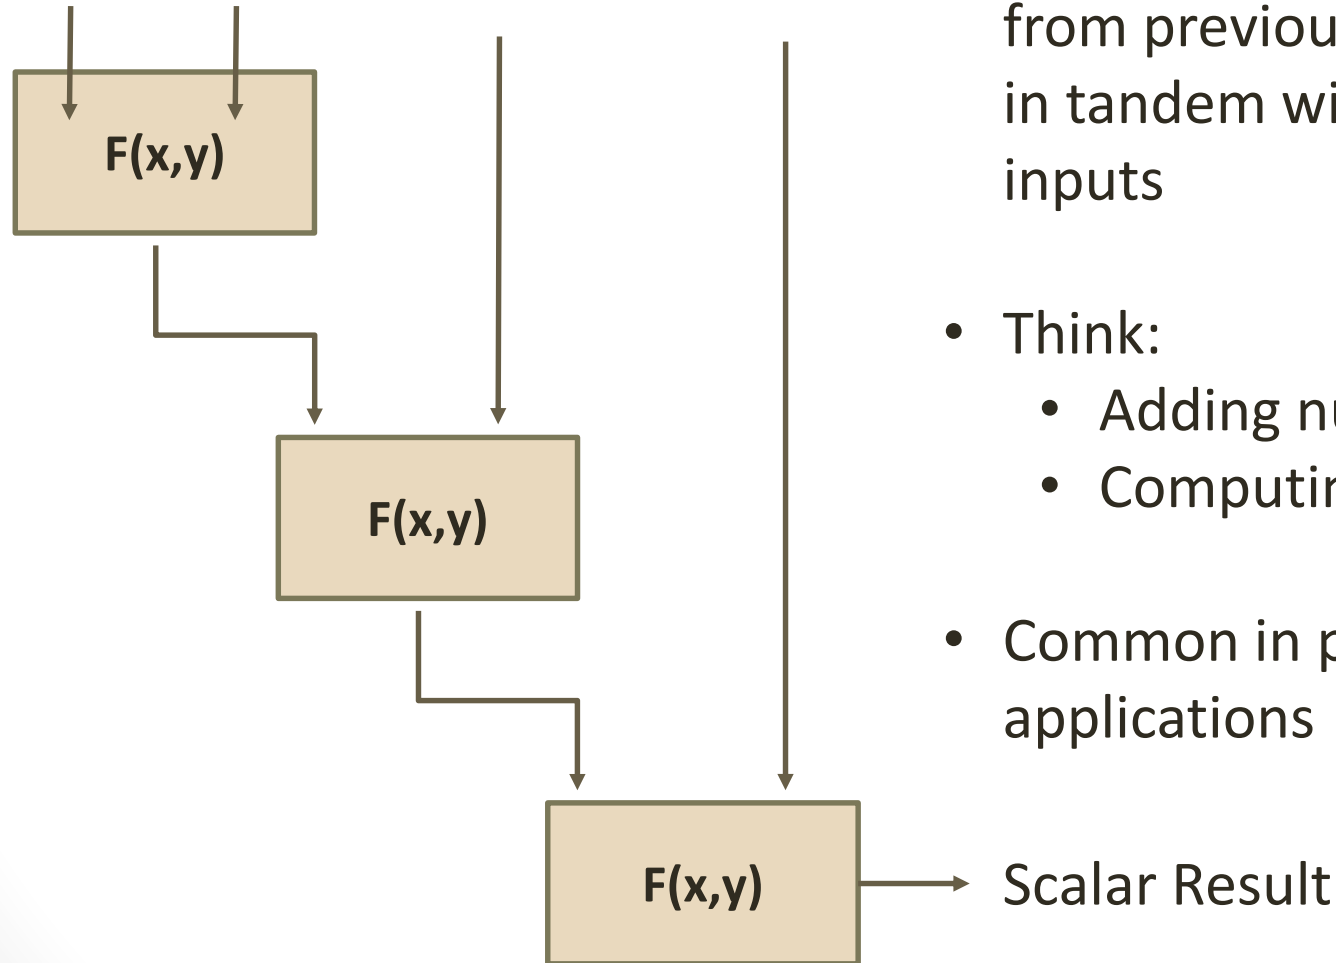
examples /

[numpy_io.py](#)

- NOTE: Arrays are ALWAYS written in Python/C-style Row-Major Order

Reduction

$A = [A[0], A[1], A[2], A[3]]$



- Sequential function evaluation using results from previous evaluations in tandem with new inputs
- Think:
 - Adding numbers
 - Computing factorials
- Common in parallel applications

Function Mapping:

- Useful shorthand for performing function evaluation on each element contained within a list of numbers
- Returns a list
- Usage:
`results = map(function_name, list_name)`

[Open this file:](#)

Parallelization Workshop / Day3-Parallel_Python /
session1_numpy / examples /
[map_reduce.py](#)

Exercise 2

Rewrite this program using

- NumPy arrays instead of lists
- array syntax instead of loops

Open this file:

Parallelization Workshop /
Day3-Parallel_Python /
session1_numpy /
exercises /
exercise2.py